

Introduction to Kernels and Regularization

M. Vazirgiannis

 DaSCIIM, LIX, École Polytechnique

October, 2021


Outline

1 Kernels

2 Support Vector Machines

3 Regularization

Mapping




- Map data points into an inner product space H with some function $\varphi : \varphi : x \rightarrow \varphi(x) \in H$
- The map φ aims to convert the nonlinear relations into linear ones.

Constructing Features

- Problems
 - Need to be an expert in the domain
 - Features may not be adequate
 - Extracting features can sometimes be computationally expensive
 - Example: second order features in 1000 dimensions.
- Solutions
 - Calculate a similarity measure in the feature space instead of the coordinates of the vectors there,
 - apply algorithms that only need the value of this measure

Kernels as distance matrices



- Define a “similarity/distance function”: $K : X \times X \rightarrow R$
- Represent a set of n data points $S = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ by the $n \times n$ matrix:
 $\mathbf{K}[ij] := K(x_i, x_j)$ where $K(x_i, x_j)$ is the distance/similarity as depicted in $K(S)$

- A kernel is a function $k : X \times X \rightarrow R$ for which the following property holds


$$k(x, x') = \langle \varphi(x), \varphi(x') \rangle$$

where φ is a mapping from X to a Hilbert (inner product) space H

$$\varphi : x \rightarrow \varphi(x) \in H$$

Kernel Example

- Quadratic Features in \mathbb{R}^2



$$\phi : x = (x_1, x_2) \rightarrow \phi(x) = (x_1^2, x_2^2, \sqrt{2}x_1x_2)$$

- Inner product in the feature space

$$\begin{aligned}\langle \phi(x), \phi(z) \rangle &= \left\langle (x_1^2, x_2^2, \sqrt{2}x_1x_2), (z_1^2, z_2^2, \sqrt{2}z_1z_2) \right\rangle \\ &= x_1^2 z_1^2 + x_2^2 z_2^2 + 2x_1 x_2 z_1 z_2 = \langle x, z \rangle^2\end{aligned}$$

- enable operation in a high-dimensional, implicit feature space
- without computing the coordinates of the data in that space - $\varphi(x)$
- simply computing inner products between the images of all pairs of data in the feature space
- need a function: $k(x, x') = \langle \varphi(x), \varphi(x') \rangle$
 - computationally cheaper than the explicit computation of the coordinates.
 - introduced for sequence data, graphs, text, images, as well as vectors.

Properties of a Kernel Matrix

- Symmetric

- K due to the symmetry of the dot product:

$$K_{ij} = K_{ji} \text{ as } \langle \phi(x), \phi(x') \rangle = \langle \phi(x'), \phi(x) \rangle$$

- K is Positive Semidefinite

- $a^T K a \geq 0$ for all $a \in \mathbb{R}^n$ and all kernel matrices $K \in \mathbb{R}^{n \times n}$.
Proof:

$$\begin{aligned}\sum_{i,j}^n a_i a_j K_{ij} &= \sum_{i,j}^n a_i a_j \langle \phi(x_i), \phi(x_j) \rangle \\ &= \left\langle \sum_i^n a_i \phi(x_i), \sum_j^n a_j \phi(x_j) \right\rangle = \left\| \sum_i^n a_i \phi(x_i) \right\|^2 \geq 0\end{aligned}$$

Constructing Kernels from Kernels

- Assuming valid kernels $k_1(x, z)$ and $k_2(x, z)$, the following are also valid kernels:
 - $k(x, z) = ck_1(x, z)$, where $c \in \mathcal{R}^+$
 - $k(x, z) = k_1(x, z) + k_2(x, z)$
 - $k(x, z) = k_1(x, z)k_2(x, z)$
 - $k(x, z) = \exp(k_1(x, z))$
 - $k(x, x') = k_a(x_a, x'_a) + k_b(x_b, x'_b)$, where $x = (x_a, x_b)$
 - $k(x, x') = k_a(x_a, x'_a)k_b(x_b, x'_b)$, where $x = (x_a, x_b)$

Typical Kernels

- Linear

$$k(x, x') = \langle x, x' \rangle$$

- Laplacian RBF

$$k(x, x') = \exp\left(-\lambda \frac{\|x - x'\|^2}{\sigma^2}\right)$$

- Gaussian RBF

$$k(x, x') = \exp\left(-\lambda \frac{\|x - x'\|^2}{\sigma^2}\right)$$


- Polynomial

$$k(x, x') = (\alpha \langle x, x' \rangle + c)^d, \alpha, c \geq 0, d \in \mathbb{N}$$


- Sigmoid

$$k(x, x') = \tanh(\alpha \langle x, x' \rangle + b)$$


Laplacian Kernel



Gaussian Kernel



Polynomial Kernel



$$k(x, x') = (\langle x, x' \rangle + c)^d, c \geq 0, d \in N$$


Outline

1 Kernels

2 Support Vector Machines

3 Regularization


- Issues that motivated SVMs:
 - bias variance tradeoff
 - Over-fitting
- For a given learning task, a finite amount of training data, the best generalization performance is achieved by jointly optimizing
 - accuracy attained on a training set,
 - “capacity”: ability to learn from any training set without error



- Goal: find a hyperplane (i.e. decision boundary) linearly separating our classes.
- Boundary equation: $\mathbf{w}^T \mathbf{x} + b = 0$
 - if $x_i : \mathbf{w}^T \mathbf{x} + b > 0$ then $y_i = 1$
 - if $x_i : \mathbf{w}^T \mathbf{x} + b < 0$ then $y_i = -1$
equivalent: $y(\mathbf{w}^T \mathbf{x} + b) \geq 1$

SVMs – distance between the boundaries

- lines are parallel, with same parameters \mathbf{w}, b
- Assume x_1 on $\mathbf{w}^T \mathbf{x} + b = 1$, the closest point of x_2 on line $\mathbf{w}^T \mathbf{x} + b = -1$. Thus $x_2 = \mathbf{w}^T \mathbf{x} + b = -1$ and $\lambda \mathbf{w}$ the distance (x_1, x_2).
- Solving for λ : $\mathbf{w}^T \mathbf{x}_2 + b = 1$ where $\mathbf{x}_2 = \mathbf{x}_1 + \lambda \mathbf{w} \Rightarrow \lambda = \frac{2}{\mathbf{w}^T \mathbf{w}} = \frac{2}{\|\mathbf{w}\|^2}$




SVMs – optimization formulation

- maximize the distance between the two boundaries defining the classes
– to avoid mis-classifications: **maximal margin**
- Objective :

$$\max \frac{2}{\sqrt{\mathbf{w}^T \mathbf{w}}} \approx \min \frac{\sqrt{\mathbf{w}^T \mathbf{w}}}{2} \approx \min \frac{\mathbf{w}^T \mathbf{w}}{2}$$

- Quadratic formulation problem:



Soft Margin extension

- We allow some miss-classification: some data points on the other side of the boundary (slack variables: $\epsilon_i > 0$ for each point x_i).
- The problem becomes:

$$\min_{\mathbf{w}, b, C} \frac{\mathbf{w}^T \mathbf{w}}{2} + C \sum_i \epsilon_i$$

Subject to: $y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \epsilon_i, \epsilon_i \geq 0 \ (\forall \mathbf{x}_i)$

- If data are not linearly separable we consider a mapping to a higher dimensional space via a function $\varphi(\mathcal{X})$. Then the optimization becomes:

$$\min_{\mathbf{w}, b, C} \frac{\mathbf{w}^T \mathbf{w}}{2} + C \sum_i \epsilon_i$$

Subject to: $y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b) \geq 1 - \epsilon_i, \epsilon_i \geq 0 \ (\forall \mathbf{x}_i)$

- Introduce Lagrangian multipliers to represent the condition
- $y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b)$ should be as close to 1 as possible :
- This condition is captured by: $\max_{\alpha_i \geq 0} \alpha_i [1 - y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b)]$
 - When $y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b) \geq 1$ the expression is maximal when $\alpha_i = 0$
 - Otherwise $y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b) < 1$, so $1 - y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b)$ is a positive value and the expression is maximal when $\alpha_i \rightarrow \infty$
- This results in penalizing (large α_i) misclassified data points, while 0 penalty to properly classified ones
- Thus we have the following formulation:

$$\min_{\mathbf{w}, b} \left[\frac{\mathbf{w}^T \mathbf{w}}{2} + \sum_i \max_{\alpha_i \geq 0} \alpha_i [1 - y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b)] \right]$$

SVMs – reformulation as a Lagrangian

$$\min_{\mathbf{w}, b} \left[\frac{\mathbf{w}^T \mathbf{w}}{2} + \sum_i \max_{\alpha_i \geq 0} \alpha_i [1 - y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b)] \right]$$

- Preventing α variables to ∞ impose constraints $0 \leq \alpha_i \leq C$
- We define the dual problem interchanging the max, min:

$$\max_{\alpha \geq 0} \left[\min_{\mathbf{w}, b} J(\mathbf{w}, b; \alpha) \right]$$

$$\text{where } J(\mathbf{w}, b; \alpha) = \frac{\mathbf{w}^T \mathbf{w}}{2} + \sum_i \alpha_i [1 - y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b)]$$

- To solve the optimization problem: $\frac{\partial J}{\partial \mathbf{w}} = 0$ hence $\mathbf{w} : \sum_i \alpha_i y_i \phi(\mathbf{x}_i)$, $\frac{\partial J}{\partial b} = 0$ hence $\sum_i \alpha_i y_i = 0$, Substitute and simplify:

$$\min_{\mathbf{w}, b} J(\mathbf{w}, b; \alpha) = \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$$

The the dual problem is:

$$\max_{\alpha \geq 0} \left[\sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j) \right]$$

Subject to: $\sum_i \alpha_i y_i = 0$ and $0 \leq \alpha_i \leq C$

- As dimensionality may be infinite computation of $\phi(\mathbf{x}_i)^T, \phi(\mathbf{x}_j)$ may be intractable

Kernel Trick:

$$K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \mathbf{x}_i^T, \mathbf{x}_j)^2 = \phi(\mathbf{x}_i)^T, \phi(\mathbf{x}_j)$$

- Thus our computation is simplified with rewriting the dual in terms of the kernel:

$$\max_{\alpha \geq 0} \left[\sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j K(\mathbf{x}_i, \mathbf{x}_j) \right]$$

- To classify a novel instance \mathbf{x} , having learned the optimal α_i parameters:

$$f(\mathbf{x}) = \text{sign} \left(\mathbf{w}^T \mathbf{x} + b \right) = \sum_i \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b$$

- Setting $\mathbf{w} = \sum_i \alpha_i y_i \phi(\mathbf{x}_i)$ and using the kernel trick
- α_i are non zero for $\varphi(x_i)$ on or close to the boundary – support vectors


Outline

1 Kernels

2 Support Vector Machines


3 Regularization

Training vs. Test error



- Model complexity causes overfitting : $(\frac{MSE_{test}}{MSE_{train}})_{degree=3} \leq (\frac{MSE_{test}}{MSE_{train}})_{degree=9}$





Polynomial Curve fitting



$$y(x, \mathbf{w}) = w_0 + w_1 x + w_2 x^2 + \dots + w_M x^M = \sum_{j=0}^M w_j x^j$$

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$


Model complexity vs. Overfitting




Polynomial coefficients

	M1	M3	M5	...	M9
W0	-3.16	-1.54	0.57		-50.04
W1		-12.13	-16.01		1618.22
W2		13.81	-16.05		-21585.65
W3			88.54		141927.36
W4			-61.42		-515139.44
W5					1083133.64
W6					-1313771.37
W7					852910.97
W8					-229423.36

Dataset size




Dataset size




Regularization

- Complex models tend to overfit
- Need to penalize the complexity of the model

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$



Regularization: Error vs. λ



	M1	M3	M5	...	M9
W0	-3.16	-1.54	0.57		-50.04
W1		-12.13	-16.01		1618.22
W2		13.81	-16.05		-21585.65
W3			88.54		141927.36
W4			-61.42		-515139.44
W5					1083133.64
W6					-1313771.37
W7					852910.97
W8					-229423.36

- A learning algorithm is stable if a small change of the input does not change the output of the algorithm much

Different risk/error types

- Empirical risk: $L_{(S)} = \frac{1}{n} \sum_{i=1}^N \delta(f(\mathbf{x}), \mathbf{y})$
- Expected risk: $L_{(D,f)} = \int_{\mathbb{X}, \mathbb{Y}} \delta(f(\mathbf{x}), \mathbf{y}) p(\mathbf{x}, \mathbf{y}) d\mathbf{x} d\mathbf{y}$

Generalization

- Empirical risk: $L_{(S)} = \frac{1}{n} \sum_{i=1}^N \delta(f(\mathbf{x}_i), \mathbf{y}_i)$
- Expected risk: $L_{(D,f)} = \int_{\mathbb{X}, \mathbb{Y}} \delta(f(\mathbf{x}), \mathbf{y}) p(\mathbf{x}, \mathbf{y}) d\mathbf{x}d\mathbf{y}$
- Generalization error: $G = L_{(D,f)} - L_{(S,f)}$
 - Difference between the *training set* and the *underlying joint probability distribution* error
 - An algorithm generalizes well if

$$\lim_{n \rightarrow \infty} G = L_{(D,f)} - L_{(S,f)} = 0$$

$p(\mathbf{x}, \mathbf{y})$ unknown probability distribution \Rightarrow impossible to compute

- Regularized Loss Minimization (RLM)
- Minimize the sum of the empirical risk + regularization function
 - measures the complexity of hypotheses/models
 - an interpretation of the regularization function is the structural risk minimization paradigm
- Regularization: stabilizer of the learning algorithm
 - stability: a slight change of its input does not change its output much

Regularized Loss Minimization (RLM)

- Learning rule jointly minimizing the empirical risk $L_s(\mathbf{w})$ and penalizing the model complexity via the regularization function $R(\mathbf{w})$

$$\arg \min_{\mathbf{w}} (L_s(\mathbf{w}) + R(\mathbf{w}))$$

- The simplest one is $R(\mathbf{w}) = \lambda \|\mathbf{w}\|^2$, λ scalar value and $\|\mathbf{w}\| = \sqrt{\sum w_i^2}$
- Therefore: $A(S) = \arg \min_{\mathbf{w}} (L_s(\mathbf{w}) + \lambda \|\mathbf{w}\|^2)$

Squared Loss – Ridge regression

- RLM rule with Tikhonov regularization to linear regression with the squared loss, we obtain the following learning rule

$$\arg \min_{\mathbf{w} \in \mathbb{R}^d} (\lambda \|\mathbf{w}\|^2 + \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (< \mathbf{w}, \mathbf{x}_i > - y_i)^2)$$

- Solving for gradient = 0 we get:

$$\mathbf{w}^{Ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

- $\lambda \geq 0$ is a parameter that controls the amount of weights' shrinkage:
 - the larger the value of λ , the greater the amount of shrinkage. The coefficients are shrunk toward zero.

Squared Loss – Ridge regression

- Equivalent formulation:

$$\arg \min_{\mathbf{w} \in \mathbb{R}^d} (\lambda \|\mathbf{w}\|^2 + \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (< \mathbf{w}, \mathbf{x}_i > - y_i)^2)$$

- subject to $\sum_{i=1}^m w_i^2 \leq t$
- We assume that mean value of y = intercept

Regularization of the intercept

- Regression $\mathbf{y} = \mathbf{w}\mathbf{X} \Rightarrow \mathbf{y} = w_0 + w_1\mathbf{x}_1 + \dots + w_m\mathbf{x}_m$
- \Rightarrow for $\mathbf{X} = 0$: $|\mathbf{y}| = w_0$
- Regularization based on the idea that overfitting on \mathbf{y} is caused by being "overly specific",
 - usually resulting in large values of \mathbf{w} elements.
- w_0 offsets the relationship: less important to the problem.
 - in case a large offset needed, regularizing it will prevent finding the correct relationship.
- $\mathbf{y} = \mathbf{w}_{1\dots M}\mathbf{X} + w_0$,
 - $\mathbf{w}_{1\dots M}$ is multiplied with the explaining variables, w_0 added to it.
- Regularizing w_0 may increase bias...

Ridge regularization computation

- $\mathbf{X}_{N \times (M+1)}$: matrix of input features (N rows, $M + 1$ columns, M weights and the intercept)
- \mathbf{y}_N : the actual outcome variable
- $\hat{\mathbf{y}}_N$: predicted values of \mathbf{y}
- \mathbf{w}_{M+1} : weights or the coefficients
- For each data point the prediction is $\hat{y}_i = \sum_{j=0}^M w_j x_{ij}$
- Cost function to be minimized: RSS (Residual Sum of Squares)

$$\text{Cost}(\mathbf{w}) = \text{RSS}(\mathbf{w}) = \sum_{i=1}^N \{y_i - \hat{y}_i\}^2 = \sum_{i=1}^N \left\{ y_i - \sum_{j=0}^M w_j x_{ij} \right\}^2$$

Ridge regularization computation

Regression algorithm – with gradient descent:

$w = 0$; η : learning rate;

while no convergence **do**

for each feature j in $\{0, 1, \dots, M\}$ **do**

 determine the gradient;

$w^{t+1} = w^t - \eta * \text{gradient};$

end

end

Gradient for the j^{th} weight:

$$\frac{\partial}{\partial w_j} = -2 \sum_{i=1}^N x_{ij} \left\{ y_i - \sum_{k=0}^M w_k x_{ik} \right\}$$

Therefore:

$$w_j^{t+1} = w_j^t + 2\eta \sum_{i=1}^N x_{ij} \left\{ y_i - \sum_{k=0}^M w_k x_{ik} \right\}$$

Ridge regularization computation

- L2 regularization loss function: $\sum_{i=1}^N \left\{ y_i - \sum_{j=0}^M w_j x_{ij} \right\}^2 + \lambda \sum_{j=0}^M w_j^2$
- The gradient:

$$\frac{\partial}{\partial w_j} Cost(\mathbf{w}) = -2 \sum_{i=1}^N x_{ij} \left\{ y_i - \sum_{k=0}^M w_k x_{ik} \right\} + 2\lambda w_j$$

$$w_j^{t+1} = w_j^t - \eta \left[-2 \sum_{i=1}^N x_{ij} \left\{ y_i - \sum_{k=0}^M w_k x_{ik} \right\} + 2\lambda w_j^t \right]$$

$$w_j^{t+1} = (1 - 2\lambda\eta) w_j^t + 2\eta \sum_{i=1}^N x_{ij} \left\{ y_i - \sum_{k=0}^M w_k x_{ik} \right\}$$

Ridge regularization computation

- Simple regression update rule:

$$w_j^{t+1} = w_j^t + 2\eta \sum_{i=1}^N x_{ij} \left\{ y_i - \sum_{k=0}^M w_k x_{ik} \right\}$$

- Ridge regression update rule:

$$w_j^{t+1} = (1 - 2\lambda\eta)w_j^t + 2\eta \sum_{i=1}^N x_{ij} \left\{ y_i - \sum_{k=0}^M w_k x_{ik} \right\}$$

- Ridge regression:
 - reduces the weights by a factor of $(1 - 2\lambda\eta)$
 - then applies the same update rule as simple linear regression
- Explains why the coefficients reduce to small numbers but never zero


Squared Loss – Ridge regression

$$A(S) = \arg \min_{\mathbf{w}} (L_S(\mathbf{w}) + \lambda ||\mathbf{w}||^2)$$

$$\mathbf{w}^{Ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

Solution is affected by the parameter λ
(shrinkage parameter)

- λ controls size of coefficients
therefore the amount of **regularization**
- for each λ value – different solution
- $\lambda \rightarrow 0$: obtain the un-regularized solution
- $\lambda \rightarrow \infty$: $\mathbf{w} = 0$
 - infinite weights lead coefficients to 0



$\lambda \rightarrow \infty$

$\lambda \rightarrow 0$


- The Lasso regression [Tibshirani, 1996] is penalizing the sum of absolute values of the weights

$$\mathbf{w}^{Lasso} = \arg \min_{\mathbf{w}} \sum_{i=1}^N (y_i - w_0 - \sum_{j=1}^m x_{ij} w_j)^2$$

- Subject to the constraint $\sum_{i=1}^m |w_i| \leq t$.
- The solution for w_0 is mean value of \mathbf{y} , thus we fit a model without the intercept w_0 .
- The constraint makes the solutions nonlinear in the y_i : no closed form expression as in ridge regression.
- Computing the lasso solution is a quadratic programming problem.

Lasso regression

- Assume $t_0 = \sum_{i=1}^m |w_i|$, w_i the weights produced by the least square solution (i. e. non regularized results)
- for values of $t \leq t_0$ ($t \geq 0$), solutions are shrunken versions of the least squares estimates
 - often, some coefficients w_j are zero
- "t" defines the number of predictors to use in a regression model
- t value to minimize expected prediction error – via cross validation



$$s = \frac{t}{\sum_{i=1}^m |w_i|}$$

- A quadratic programming problem; can be tackled by numerical analysis algorithms.

Incremental Forward Stepwise (FS)¹ Regression:

- 1 Start with residual $\mathbf{r} = \mathbf{y} - \bar{\mathbf{y}}$ and all $w_j = 0$;
repeat
 - 2 Find predictor \mathbf{x}_j most correlated with \mathbf{r} ;
 - 3 Update $w_j \leftarrow w_j + \delta_j$, where $\delta_j = \epsilon * \text{sign}[\text{corr}(\mathbf{r}, \mathbf{x}_j)]$;
 - 4 Update $\mathbf{r} \leftarrow \mathbf{r} - \delta_j \mathbf{x}_j$
- until** no predictor has any correlation with \mathbf{r} ;

¹Forward stagewise regression and the monotone lasso, Trevor Hastie, Jonathan Taylor, Robert Tibshirani, Guenther Walther, *Electronic Journal of Statistics*, Vol. 1 (2007) 1–29
ISSN: 1935-7524. <https://arxiv.org/pdf/0705.0269.pdf>

Computation of the Lasso solutions

- Least angle regression procedure is a better approach
 - exploits the special structure of the lasso problem,
 - efficient way to compute the solutions simultaneously for all the values of " w_i "
 - least angle is based on the forward stepwise regression

Least angle regression algorithm:

1 Set all coefficients $w_j = 0$;

repeat

2 Find the predictor \mathbf{x}_j most correlated with \mathbf{y} ;

3 Increase the coefficient w_j in the direction of the sign of its correlation with \mathbf{y} ;

4 Take residuals $\mathbf{r} = \mathbf{y} - \bar{\mathbf{y}}$ along the way;

5 Stop when some other predictor \mathbf{x}_k has as much correlation with \mathbf{r} as \mathbf{x}_j has;

6 Increase $(w_j; w_k)$ in their joint least squares direction;

7 until some other predictor \mathbf{x}_m has as much correlation with the residual \mathbf{r} ;

until all predictors are in the model;

- After p steps w_i reach their ridge (L2) solution.

Least angle regression algorithm:

- 1 Set all coefficients $w_j = 0$;
 - 2 repeat
 - 3 Find the predictor \mathbf{x}_j most correlated with \mathbf{y} ;
 - 4 Increase the coefficient w_j in the direction of the sign of its correlation with \mathbf{y} ;
 - 5 Take residuals $\mathbf{r} = \mathbf{y} - \bar{\mathbf{y}}$ along the way;
 - 6 Stop when some other predictor \mathbf{x}_k has as much correlation with \mathbf{r} as \mathbf{x}_j has;
 - 7 Increase $(w_j w_k)$ in their joint least squares direction;
 - 8 until some other predictor \mathbf{x}_m has as much correlation with the residual \mathbf{r} ;
- until all predictors are in the model;**


- **LARS: Lasso modification:**

If a non-zero coefficient hits zero, drop it from the active set and recompute the current joint least squares direction.



Comparing ridge and lasso

Contours of the error and constraint functions:

- solid blue areas – constraint regions: $|\beta_1| + |\beta_2| \leq t$ and $\beta_1^2 + \beta_2^2 \leq t^2$.
- red ellipses: contours of the least squares error function.



Comparing ridge and lasso




2

- variable coefficient estimates: Ridge regression (*left*), LASSO (*right*) for the red wine data plotted versus $\log \lambda$.
- Upper part of the plot shows the number of non-zero coefficients in the regression model for a given $\log(\lambda)$.

²L.E. Melkumova et al. / Procedia Engineering 201 (2017) 746–755

Ridge regularization effect on regression weights

- Regression: RSS vs. model complexity (polynomial degree)



<https://www.analyticsvidhya.com/blog/2016/01/complete-tutorial-ridge-lasso-regression-python/>

Ridge regularization effect on regression weights


- Regression: coefficient values vs. model complexity (polynomial degree)

	rss	intercept	coef_x_1	coef_x_2	coef_x_3	coef_x_4	coef_x_5	coef_x_6	coef_x_7	coef_x_8	coef_x_9	coef_x_10	coef_x_11	c
model_pow_1	3.3	2	-0.62	NaN	NaN	1								
model_pow_2	3.3	1.9	-0.58	-0.006	NaN	NaN	1							
model_pow_3	1.1	-1.1	3	-1.3	0.14	NaN	NaN	1						
model_pow_4	1.1	-0.27	1.7	-0.53	-0.036	0.014	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1
model_pow_5	1	3	-5.1	4.7	-1.9	0.33	-0.021	NaN	NaN	NaN	NaN	NaN	NaN	1
model_pow_6	0.99	-2.8	9.5	-9.7	5.2	-1.6	0.23	-0.014	NaN	NaN	NaN	NaN	NaN	1
model_pow_7	0.93	19	-56	69	-45	17	-3.5	0.4	-0.019	NaN	NaN	NaN	NaN	1
model_pow_8	0.92	43	-1.4e+02	1.8e+02	-1.3e+02	58	-15	2.4	-0.21	0.0077	NaN	NaN	NaN	1
model_pow_9	0.87	1.7e+02	-6.1e+02	9.6e+02	-8.5e+02	4.6e+02	-1.6e+02	37	-5.2	0.42	-0.015	NaN	NaN	1
model_pow_10	0.87	1.4e+02	-4.9e+02	7.3e+02	-6e+02	2.9e+02	-87	15	-0.81	-0.14	0.026	-0.0013	NaN	1
model_pow_11	0.87	-75	5.1e+02	-1.3e+03	1.9e+03	-1.6e+03	9.1e+02	-3.5e+02	91	-16	1.8	-0.12	0.0034	1
model_pow_12	0.87	-3.4e+02	1.9e+03	-4.4e+03	6e+03	-5.2e+03	3.1e+03	-1.3e+03	3.8e+02	-80	12	-1.1	0.062	-1
model_pow_13	0.86	3.2e+03	-1.8e+04	4.5e+04	-6.7e+04	6.6e+04	-4.6e+04	2.3e+04	-8.5e+03	2.3e+03	-4.5e+02	62	-5.7	0
model_pow_14	0.79	2.4e+04	-1.4e+05	3.8e+05	-6.1e+05	6.6e+05	-5e+05	2.8e+05	-1.2e+05	3.7e+04	-8.5e+03	1.5e+03	-1.8e+02	1
model_pow_15	0.7	-3.6e+04	2.4e+05	-7.5e+05	1.4e+06	-1.7e+06	1.5e+06	-1e+06	5e+05	-1.9e+05	5.4e+04	-1.2e+04	1.9e+03	-1

<https://www.analyticsvidhya.com/blog/2016/01/complete-tutorial-ridge-lasso-regression-python/>

Ridge regularization effect on regression weights

- Ridge Regression: RSS vs. λ parameter – shrinkage



<https://www.analyticsvidhya.com/blog/2016/01/complete-tutorial-ridge-lasso-regression-python/>


Ridge regularization effect on regression weights

- Ridge Regression: coefficient values (regression weights) vs. λ (α in table) parameter shrinkage

	rss	intercept	coef_x_1	coef_x_2	coef_x_3	coef_x_4	coef_x_5	coef_x_6	coef_x_7	coef_x_8	coef_x_9	coef_x_10	coef_x_11	c
alpha_1e-15	0.87	95	-3e+02	3.8e+02	-2.4e+02	66	0.96	-4.8	0.64	0.15	-0.026	-0.0054	0.00086	0
alpha_1e-10	0.92	11	-29	31	-15	2.9	0.17	-0.091	0.002	0.00064	2.4e-05	-2e-05	-	-
alpha_1e-08	0.95	1.3	-1.5	1.7	-0.68	0.039	0.016	0.00016	-0.00036	-5.4e-05	2.9e-07	1.1e-06	1.9e-07	2
alpha_0.0001	0.96	0.56	0.55	-0.13	-0.026	-0.0028	-0.00011	4.1e-05	1.5e-05	3.7e-06	7.4e-07	1.3e-07	1.9e-08	1
alpha_.001	1	0.82	0.31	-0.087	-0.02	-0.0028	-0.00022	1.8e-05	1.2e-05	3.4e-06	7.3e-07	1.3e-07	1.9e-08	1
alpha_.01	1.4	1.3	-0.088	-0.052	-0.01	-0.0014	-0.00013	7.2e-07	4.1e-06	1.3e-06	3e-07	5.6e-08	9e-09	1
alpha_1	5.6	0.97	-0.14	-0.019	-0.003	-0.00047	-7e-05	-9.9e-06	-1.3e-06	-1.4e-07	-9.3e-09	1.3e-09	7.8e-10	2
alpha_5	14	0.55	-0.059	-0.0085	-0.0014	-0.00024	-4.1e-05	-6.9e-06	-1.1e-06	-1.9e-07	-3.1e-08	-5.1e-09	-8.2e-10	-
alpha_10	18	0.4	-0.037	-0.0055	-0.00095	-0.00017	-3e-05	-5.2e-06	-9.2e-07	-1.6e-07	-2.9e-08	-5.1e-09	-9.1e-10	-
alpha_20	23	0.28	-0.022	-0.0034	-0.0006	-0.00011	-2e-05	-3.6e-06	-6.6e-07	-1.2e-07	-2.2e-08	-4e-09	-7.5e-10	-

- 1 The RSS increases with λ , while model complexity reduces
- 2 Even a **small** λ (e^{-15}) results in significant reduction coefficients magnitude. Compare the coefficients first row to the last row of simple linear regression table.
- 3 High λ values lead to significant under-fitting
 - Note rapid increase in RSS for $\lambda > 1$
 - The coefficients are **very very small**, they are **NOT zero**

Lasso regression effect on regression weights



Model complexity decreases with increase in the values of λ .

- Notice straight line at $\lambda = 1$

<https://www.analyticsvidhya.com/blog/2016/01/complete-tutorial-ridge-lasso-regression-python/>

Regularization and sparsity

- For the same values of λ (α in table), lasso coefficients \ll ridge coefficients.
- For the same λ , lasso has higher RSS (poorer fit) as compared to ridge regression.
- Many of the coefficients are zero even for very small values of alpha.
- Inferences #1, 2 might not generalize always but will hold for many cases. The real difference from ridge is coming out in the last inference.

	rss	intercept	coef_x_1	coef_x_2	coef_x_3	coef_x_4	coef_x_5	coef_x_6	coef_x_7	coef_x_8	coef_x_9	coef_x_10	coef_x_11	co
alpha_1e-15	0.96	0.22	1.1	-0.37	0.00089	0.0016	-0.00012	-6.4e-05	-6.3e-06	1.4e-06	7.8e-07	2.1e-07	4e-08	5.4
alpha_1e-10	0.96	0.22	1.1	-0.37	0.00088	0.0016	-0.00012	-6.4e-05	-6.3e-06	1.4e-06	7.8e-07	2.1e-07	4e-08	5.4
alpha_1e-08	0.96	0.22	1.1	-0.37	0.00077	0.0016	-0.00011	-6.4e-05	-6.3e-06	1.4e-06	7.8e-07	2.1e-07	4e-08	5.3
alpha_1e-05	0.96	0.5	0.6	-0.13	-0.038	-0	0	0	0	7.7e-06	1e-06	7.7e-08	0	0
alpha_0.0001	1	0.9	0.17	-0	-0.048	-0	-0	0	0	9.5e-06	5.1e-07	0	0	0
alpha_0.001	1.7	1.3	-0	-0.13	-0	-0	-0	0	0	0	0	0	1.5e-08	7.5
alpha_0.01	3.6	1.8	-0.55	-0.00056	-0	-0	-0	-0	-0	-0	-0	0	0	0
alpha_1	37	0.038	-0	-0	-0	-0	-0	-0	-0	-0	-0	-0	-0	-0
alpha_5	37	0.038	-0	-0	-0	-0	-0	-0	-0	-0	-0	-0	-0	-0
alpha_10	37	0.038	-0	-0	-0	-0	-0	-0	-0	-0	-0	-0	-0	-0

HIGH SPARSITY

<https://www.analyticsvidhya.com/blog/2016/01/complete-tutorial-ridge-lasso-regression-python/>

Elastic Net

- $p > n$ (#variables greater than sample size), lasso can select only n variables
 - select only one from any set of highly correlated variables
 - even when $n > p$, if the variables are strongly correlated, ridge regression tends to perform better
-
- Combines L1 and L2 regularization to balance out the pros and cons of ridge and lasso regression

$$\arg \min_{\mathbf{w} \in \mathbb{R}^p} \{|\mathbf{y} - \mathbf{X}\mathbf{w}|^2 + \lambda_1 |\mathbf{w}| + \lambda_2 |\mathbf{w}|^2\}$$

- Subject to: $(1 - \alpha)|\mathbf{w}| + \alpha|\mathbf{w}|^2 \leq t$ with $\alpha = \frac{\lambda_2}{\lambda_1 + \lambda_2}$

Other regularization approaches

Group Lasso

- Regularization based on groups of variables
- Lasso selects one of them – we might need all variables of a group (i.e. topic modelling, biological applications)

Objective function:

$$\arg \min_{\mathbf{w} \in \mathbb{R}^p} \left\{ |y - \sum_{j=1}^J \mathbf{X}_j w_j|^2 + \lambda \sum_{j=1}^J |w_j|_{\mathbf{K}_j} \right\}$$

where J : the variable groups and $|w_j|_{\mathbf{K}_j} = (\mathbf{w}^T \mathbf{K}_j \mathbf{w})^{\frac{1}{2}}$.

Properties of Group Lasso:

- ① entire groups are eliminated
- ② all variables of a group are present with the same weight
- ③ if groups contain 1 variable \Rightarrow lasso

To weight variables within groups: Sparse Group Lasso

Regularization – Conclusion

- *Ridge*: all of the features retained in the model. Major advantage: coefficient shrinkage thus reducing model complexity.
- *Lasso*: Along with shrinking coefficients, lasso performs feature selection as some of the coefficients become exactly zero.

Typical Use Cases

- *Ridge*: prevent overfitting – includes all features, useful in case of very high number of features.
- *Lasso*: provides sparse solutions – suitable for cases where # features is large: great computational advantage as features with zero coefficients can be ignored (feature selection).

Presence of Highly Correlated Features

- *Ridge*: works well for highly correlated features – includes all of them in the model, coefficients will adjust.
- *Lasso*: selects feature among the highly correlated ones, reduces the coefficients of the rest to zero.
For highly correlated variables even small λ values give significant sparsity.

References I

-  Christopher M. Bishop
Pattern Recognition and Machine Learning
2006
-  Trevor Hastie, Robert Tibshirani, Jerome Friedman
The Elements of Statistical Learning: Data Mining, Inference, and Prediction
Second Edition. 2009
-  Shai Shalev-Shwartz, Shai Ben-David
Understanding Machine Learning: theory and algorithms
Cambridge University Press. 2014
-  Robert Tibshirani
Regression shrinkage and selection via the Lasso
Journal of the Royal Statistics Society. 58(1), 267-288. 1996
<http://statweb.stanford.edu/~tibs/lasso/lasso.pdf>
-  <http://people.inf.elte.hu/kiss/13dwhdm/roc.pdf>

References II



Hui Zou, Trevor Hastie

Regularization and Variable Selection via the Elastic Net

Journal of the Royal Statistical Society. Series B (Statistical Methodology). Wiley.67(2): 301–20. 2005



Jerome Friedman, Trevor Hastie, Robert Tibshirani

A note on the group lasso and a sparse group lasso

<https://arxiv.org/pdf/1001.0736.pdf>

Acknowledgments



Dr. Nikos Tziortziotis
for his contributions in the Kernels section